

Logistic Regression - Heart Data

import libraries

```
In [65]: import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression

from sklearn.metrics import confusion_matrix, accuracy_score, classification_report, roc_curve, roc_auc_score

import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import warnings
warnings.filterwarnings("ignore")
```

Data Gathering

```
In [2]: df = pd.read_csv("heart.csv")
df
```

```
Out[2]:   age  gender  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slope  ca  thal  target
      0    63       1     3      145    233     1      0     150      0      2.3      0     0     1      1
      1    37       1     2      130    250     0      1     187      0      3.5      0     0     2      1
      2    41       0     1      130    204     0      0     172      0      1.4      2     0     2      1
      3    56       1     1      120    236     0      1     178      0      0.8      2     0     2      1
      4    57       0     0      120    354     0      1     163      1      0.6      2     0     2      1
     ...
  298    57       0     0      140    241     0      1     123      1      0.2      1     0     3      0
  299    45       1     3      110    264     0      1     132      0      1.2      1     0     3      0
  300    68       1     0      144    193     1      1     141      0      3.4      1     2     3      0
  301    57       1     0      130    131     0      1     115      1      1.2      1     1     3      0
  302    57       0     1      130    236     0      0     174      0      0.0      1     1     2      0
```

303 rows × 14 columns

EDA

```
In [3]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   age         303 non-null    int64  
 1   gender       303 non-null    int64  
 2   cp          303 non-null    int64  
 3   trestbps    303 non-null    int64  
 4   chol         303 non-null    int64  
 5   fbs          303 non-null    int64  
 6   restecg     303 non-null    int64  
 7   thalach      303 non-null    int64  
 8   exang        303 non-null    int64  
 9   oldpeak      303 non-null    float64 
 10  slope        303 non-null    int64  
 11  ca           303 non-null    int64  
 12  thal         303 non-null    int64  
 13  target        303 non-null    int64  
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
In [4]: df['target'].value_counts()
```

```
Out[4]: target
1    165
0    138
Name: count, dtype: int64
```

```
In [5]: 165/303
```

```
Out[5]: 0.5445544554455446
```

```
In [ ]:
```

Feature Selection

```
In [ ]: 1. Filter Method:
```

```
2. Wrapper Method
```

```
3. Embedded Method:
```

```
    1. Lasso Regression
```

Model Training

Train Test Split

```
In [6]: x = df.drop('target', axis = 1)
y = df['target']
# y
```

```
In [7]: x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.2, random_state=10, stratify = y)
y_train.value_counts()
```

```
Out[7]: target
1    132
0    110
Name: count, dtype: int64
```

```
In [8]: 132/242
```

```
Out[8]: 0.5454545454545454
```

```
In [9]: # x_train
```

```
In [10]: # x_test
```

```
In [11]: log_clf = LogisticRegression()
log_clf.fit(x_train, y_train)
```

```
Out[11]: ▾ LogisticRegression ⓘ ?
```

```
LogisticRegression()
```

Evaluation

```
In [12]: y_pred_test_prob = log_clf.predict_proba(x_test)
y_pred_test_prob[:6]

# [0.07054224, 0.92945776], >> Class1
# [0.26737039, 0.73262961], >> Class1
# [0.82159417, 0.17840583], >> Class0
# [0.98835195, 0.01164805], >> Class0
# [0.84237389, 0.15762611], >> Class0
# [0.01919849, 0.98080151], >> Class1
```

```
Out[12]: array([[0.07054224, 0.92945776],
[0.26737039, 0.73262961],
[0.82159417, 0.17840583],
[0.98835195, 0.01164805],
[0.84237389, 0.15762611],
[0.01919849, 0.98080151]])
```

```
In [13]: y_pred_test = log_clf.predict(x_test)
y_pred_test[50:55]
```

```
Out[13]: array([0, 1, 1, 0, 1], dtype=int64)
```

```
In [14]: y_test[50:55]
```

```
Out[14]: 165    0
149    1
278    0
298    0
90     1
Name: target, dtype: int64
```

```
In [15]: confusion_matrix(y_test, y_pred_test)
```

```
Out[15]: array([[23,  5],
   [ 3, 30]], dtype=int64)

In [16]: y_test.value_counts()

Out[16]: target
1    33
0    28
Name: count, dtype: int64

In [17]: y_test
```

```
Out[17]: 94    1
99    1
280   0
174   0
176   0
..
115   1
101   1
219   0
253   0
243   0
Name: target, Length: 61, dtype: int64
```

```
In [18]: (23+30)/61

Out[18]: 0.8688524590163934
```

```
In [19]: accuracy_score(y_test, y_pred_test)
```

```
Out[19]: 0.8688524590163934
```

```
In [20]: log_clf.score(x_test, y_test)
```

```
Out[20]: 0.8688524590163934
```

Training Data Evaluation

```
In [21]: y_pred_train = log_clf.predict(x_train)
y_pred_train[60:65]
```

```
Out[21]: array([0, 1, 1, 1, 1], dtype=int64)
```

```
In [22]: y_train[60:65]
```

```
Out[22]: 247    0
37     1
293   0
114   1
19    1
Name: target, dtype: int64
```

```
In [23]: log_clf.predict_proba(x_train)[60:65]
```

```
Out[23]: array([[0.78384075,  0.21615925],
   [0.42148032,  0.57851968],
   [0.4989611 ,  0.5010389 ],
   [0.11473814,  0.88526186],
   [0.0618728 ,  0.9381272 ]])
```

```
In [24]: confusion_matrix(y_train, y_pred_train)
```

```
Out[24]: array([[ 83,  27],
   [ 10, 122]], dtype=int64)
```

```
In [25]: (83+122)/242
```

```
Out[25]: 0.8471074380165289
```

```
In [26]: 122/132
```

```
Out[26]: 0.9242424242424242
```

```
In [27]: TP = 122
FP = 27
Precision = TP/(TP+FP)
Precision
```

```
Out[27]: 0.8187919463087249
```

```
In [55]: clf_report = classification_report(y_train, y_pred_train)
print("Classification Report :\n", clf_report)
```

```
Classification Report :
    precision    recall    f1-score   support
    0            0.89     0.75     0.82      110
    1            0.82     0.92     0.87      132

    accuracy                           0.85      242
   macro avg       0.86     0.84     0.84      242
weighted avg       0.85     0.85     0.85      242
```

```
In [29]: clf_report = classification_report(y_test, y_pred_test)
print("Classification Report :\n", clf_report)
```

```
Classification Report :
    precision    recall    f1-score   support
    0            0.88     0.82     0.85      28
    1            0.86     0.91     0.88      33

    accuracy                           0.87      61
   macro avg       0.87     0.87     0.87      61
weighted avg       0.87     0.87     0.87      61
```

```
In [ ]:
```

```
In [42]: y_pred_train_prob = log_clf.predict_proba(x_train)
fpr, tpr, thresh = roc_curve(y_train, y_pred_train_prob[:,1])
```

```
In [44]: fpr
```

```
Out[44]: array([0.          , 0.          , 0.          , 0.00909091, 0.00909091,
 0.01818182, 0.01818182, 0.02727273, 0.02727273, 0.03636364,
 0.03636364, 0.04545455, 0.04545455, 0.05454545, 0.05454545,
 0.06363636, 0.06363636, 0.07272727, 0.07272727, 0.08181818,
 0.08181818, 0.09090909, 0.09090909, 0.09090909, 0.09090909,
 0.1        , 0.1        , 0.11818182, 0.11818182, 0.12727273,
 0.12727273, 0.13636364, 0.13636364, 0.14545455, 0.14545455,
 0.16363636, 0.16363636, 0.17272727, 0.17272727, 0.21818182,
 0.21818182, 0.22727273, 0.22727273, 0.23636364, 0.23636364,
 0.26363636, 0.26363636, 0.28181818, 0.28181818, 0.30909091,
 0.30909091, 0.37272727, 0.37272727, 0.39090909, 0.39090909,
 0.42727273, 0.42727273, 0.47272727, 0.47272727, 0.5        ,
 0.5        , 0.64545455, 0.64545455, 1.          ])
```

```
In [46]: tpr
```

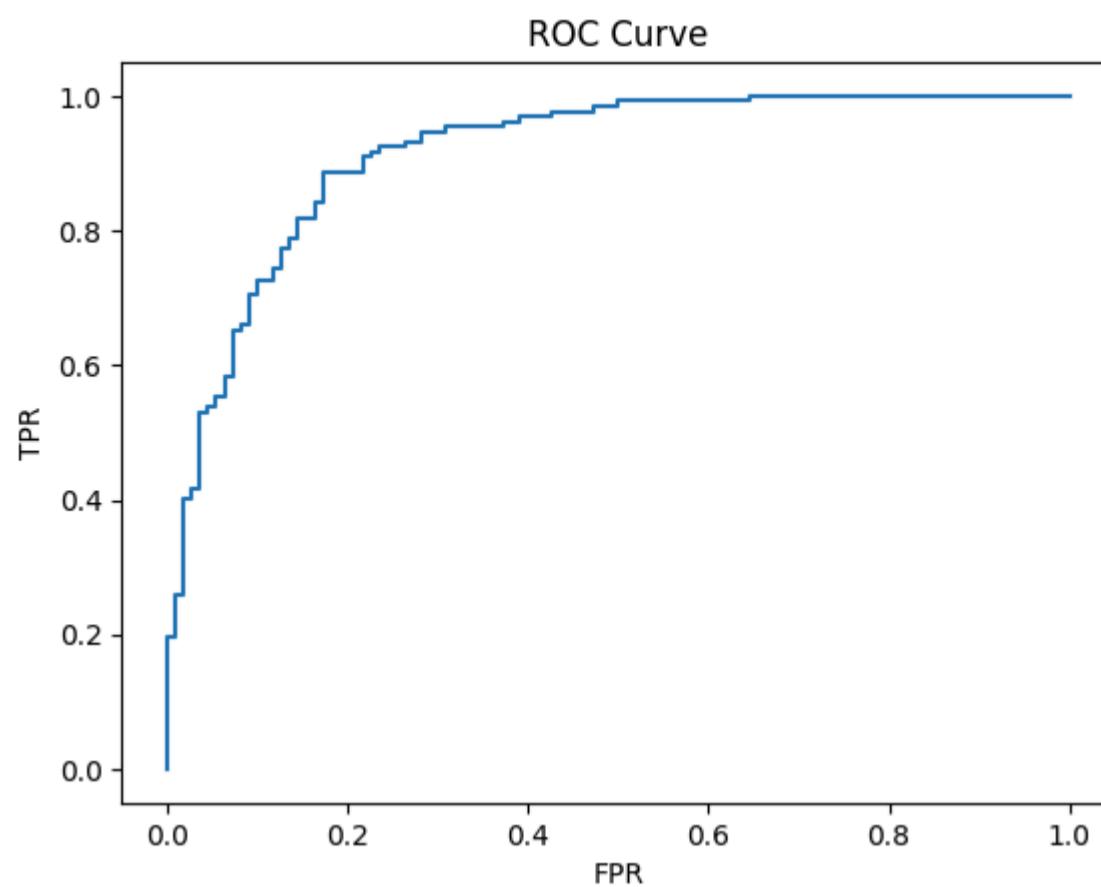
```
Out[46]: array([0.          , 0.00757576, 0.1969697 , 0.1969697 , 0.25757576,
 0.25757576, 0.40151515, 0.40151515, 0.41666667, 0.41666667,
 0.53030303, 0.53030303, 0.53787879, 0.53787879, 0.5530303 ,
 0.5530303 , 0.58333333, 0.58333333, 0.65151515, 0.65151515,
 0.65909091, 0.65909091, 0.68181818, 0.6969697 , 0.70454545,
 0.70454545, 0.72727273, 0.72727273, 0.74242424, 0.74242424,
 0.77272727, 0.77272727, 0.78787879, 0.78787879, 0.81818182,
 0.81818182, 0.84090909, 0.84090909, 0.88636364, 0.88636364,
 0.90909091, 0.90909091, 0.91666667, 0.91666667, 0.92424242,
 0.92424242, 0.93181818, 0.93181818, 0.9469697 , 0.9469697 ,
 0.95454545, 0.95454545, 0.96212121, 0.96212121, 0.96969697,
 0.96969697, 0.97727273, 0.97727273, 0.98484848, 0.98484848,
 0.99242424, 0.99242424, 1.          , 1.          ])
```

```
In [48]: thresh
```

```
Out[48]: array([      inf, 0.99448755, 0.95321669, 0.95289286, 0.9378883 ,
 0.93209381, 0.88684339, 0.88657238, 0.88466559, 0.88273015,
 0.8191005 , 0.81674828, 0.81637076, 0.81158425, 0.8095993 ,
 0.80711642, 0.7850566 , 0.78488461, 0.76050925, 0.75211307,
 0.74879297, 0.74767768, 0.73921676, 0.73666853, 0.73457792,
 0.72893221, 0.70343623, 0.70122651, 0.68876316, 0.68636715,
 0.66671936, 0.6630252 , 0.6459502 , 0.64170872, 0.62897988,
 0.61801926, 0.60482762, 0.6046812 , 0.53592203, 0.52755954,
 0.51433146, 0.50748671, 0.50261164, 0.50184913, 0.50122575,
 0.49420772, 0.48432043, 0.45922821, 0.4445082 , 0.4118362 ,
 0.40306403, 0.25813904, 0.2455269 , 0.24301206, 0.22919007,
 0.21615925, 0.21538947, 0.14613663, 0.13687115, 0.11101719,
 0.10812541, 0.05359264, 0.05058366, 0.00170894])
```

```
In [52]: plt.plot(fpr,tpr)
plt.xlabel("FPR")
plt.ylabel("TPR")
plt.title("ROC Curve")
```

```
Out[52]: Text(0.5, 1.0, 'ROC Curve')
```



```
In [54]: px.line(x = fpr, y = tpr, labels = {'x' : "FPR", 'y': "TPR"})
```

```
In [57]: clf_report = classification_report(y_train, y_pred_train)
print("Classification Report :\n", clf_report)
```

	precision	recall	f1-score	support
0	0.89	0.75	0.82	110
1	0.82	0.92	0.87	132
accuracy			0.85	242
macro avg	0.86	0.84	0.84	242
weighted avg	0.85	0.85	0.85	242

```
In [59]: np.where(tpr >= 0.94)
```

```
Out[59]: (array([48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59, 59, 60, 60, 61, 62, 63],
      dtype=int64),)
```

```
In [61]: tpr[48]
```

```
Out[61]: 0.9469696969696969
```

```
In [63]: thresh[48]
```

```
Out[63]: 0.4445081983336496
```

```
In [64]: fpr[48]
```

```
Out[64]: 0.2818181818181818
```

```
In [67]: roc_auc_score(y_train, y_pred_train_prob[:,1])
```

```
Out[67]: 0.9151515151515152
```

```
In [68]: prob = y_pred_train_prob[:,1]
```

```
In [70]: logit_values = np.log(prob/(1-prob))
```

```
In [ ]:
```